

SENIOR SDET & LEAD QA

21 Interview Questions • Detailed Answers • Architecture • Examples • Follow-ups

A practical preparation guide for explaining engineering decisions, not just listing tools.

How to use this guide

For each question, learn the core concept, practise the example, then explain the trade-off and failure modes. Senior interviews assess whether you can choose a solution, justify it, measure its effect and lead the team through delivery.

Technology baseline: Selenium Grid 4, Playwright with TypeScript, REST APIs, Jenkins/CI, containers, observability and modern performance testing. Verify version-specific commands against the versions used by the target company.

Contents

- Architecture & Framework Design (1–4)
- API & Mocking (5–10)
- Performance & Security (11–12)
- Observability & Leadership (13–17)
- Testing Strategy (18–19)
- Coding / DSA (20–21)
- Senior-level closing checklist and source references

A strong senior-level answer usually follows this pattern:

Step	What to explain
Context	What problem are we solving and what constraints matter?
Design	What components, boundaries and data flows are involved?
Trade-offs	Why this approach rather than alternatives? What are the risks?
Evidence	What metrics, logs, tests or outcomes prove it works?
Ownership	How do you roll it out, support it and improve it?


Part 1 — Architecture & Framework Design

1. How does Selenium Grid 4 work? Explain its asynchronous architecture.

Interview-ready explanation

Selenium Grid 4 distributes WebDriver sessions across multiple machines or containers. The client sends WebDriver commands to the Grid Router. A new-session request is placed in the New Session Queue. The Distributor matches requested capabilities (browser, platform, version) to an available Node slot, asks that Node to create the session, and the session ID-to-Node mapping is stored in the Session Map. For later commands, the Router looks up the session and forwards the command to the correct Node. The Node controls the browser and returns the response. The phrase “asynchronous architecture” needs precision: Grid uses both synchronous HTTP request/response communication and asynchronous Event Bus messages. The Event Bus decouples internal events such as Node registration and status/heartbeat notifications; most WebDriver commands still require synchronous responses.

Practical example



```
# Start a simple Hub and Node (illustrative Selenium 4 commands)
java -jar selenium-server-<version>.jar hub
java -jar selenium-server-<version>.jar node --hub http://<hub-host>:4444
```

Core components: Router = entry point; New Session Queue = waiting room; Distributor = schedules sessions; Node = executes browser sessions; Session Map = session routing lookup; Event Bus = asynchronous internal event channel.

Senior-level considerations

Do not say all Grid communication is asynchronous. Explain the distinction between synchronous WebDriver command/response traffic and asynchronous internal events. Discuss queue timeout, slot capacity, browser/driver compatibility, node draining, observability and isolation. Cite the official Grid architecture documentation in your preparation notes.

Likely follow-up questions


- Why does Grid need both HTTP and Event Bus?
- What happens when all Node slots are busy?
- How would you scale Nodes in Kubernetes?
- How do you investigate a session-creation timeout?

2. What does an ideal automation framework look like?

Interview-ready explanation

An ideal framework is maintainable, observable, scalable and aligned with product risk. It should make tests easy to write while making failures easy to diagnose. I design it in layers: test/specification layer, business-flow or service layer, page/API client layer, test-data and fixtures layer, configuration/environment layer, and reporting/observability layer. Cross-cutting capabilities include logging, screenshots/traces, retries with clear policy, tagging, parallel execution, secrets management and CI integration.

Practical example



```
repo/
  tests/           # scenarios and assertions
  features/        # Gherkin (if BDD is used)
  pages/           # UI page objects
  api/             # clients, schemas, contracts
  fixtures/        # setup, teardown, auth
  data/            # synthetic test data
  config/          # environment configuration
  utils/           # logging, waits, helpers
  reports/         # generated artifacts
```

Senior-level considerations

Keep business intent separate from mechanics. Prefer stable test IDs and semantic locators; avoid brittle XPath and hard waits. Do not hide product defects inside broad try/catch blocks. Use explicit ownership, code review, versioned dependencies, examples and migration guides. A framework is a product: measure adoption, reliability, execution time and maintenance cost.

Likely follow-up questions

- How do you prevent a framework becoming a collection of utility methods?
- Where should assertions live?
- How do you manage environment-specific configuration?
- How do you decide whether to use Page Object Model or Screenplay?

3. How do you maintain a framework for parallel execution (Selenium + Playwright)?

Interview-ready explanation

Parallel safety starts with isolation. Each test should own its browser context/session, test data, account or tenant, downloads and output paths. Avoid mutable static WebDriver instances, shared test accounts with conflicting state, fixed ports and shared files. Selenium can use ThreadLocal driver management when using TestNG/JUnit parallel execution, with reliable teardown. Playwright's test runner provides isolated browser contexts per test by default; use fixtures and workers deliberately. At CI level, shard suites, cap concurrency according to Grid slots and backend capacity, and preserve traces/logs per test.

Practical example

3. How do you maintain a framework for parallel execution (Selenium + Playwright)?

Interview-ready explanation

Parallel safety starts with isolation. Each test should own its browser context/session, test data, account or tenant, downloads and output paths. Avoid mutable static WebDriver instances, shared test accounts with conflicting state, fixed ports and shared files. Selenium can use ThreadLocal driver management when using TestNG/JUnit parallel execution, with reliable teardown. Playwright's test runner provides isolated browser contexts per test by default; use fixtures and workers deliberately. At CI level, shard suites, cap concurrency according to Grid slots and backend capacity, and preserve traces/logs per test.

Practical example

```
# Selenium Java: simplified ThreadLocal driver pattern
public final class DriverManager {
    private static final ThreadLocal<WebDriver> DRIVER = new ThreadLocal<>();

    public static WebDriver getDriver() {
        return DRIVER.get();
    }

    public static void setDriver(WebDriver driver) {
        DRIVER.set(driver);
    }

    public static void quitDriver() {
        WebDriver driver = DRIVER.get();
        if (driver != null) {
            driver.quit();
            DRIVER.remove();
        }
    }
}
```

Senior-level considerations

Set concurrency from the narrowest bottleneck: Grid slots, browser memory/CPU, API rate limits, database capacity and test-data constraints. Use unique run IDs and deterministic cleanup. Retries should be diagnostic, not a way to turn red builds green silently. Track first-run pass rate separately from retry-recovered pass rate.

Likely follow-up questions


- How do you choose worker count?
- How do you avoid test-data collisions?
- When should a test be serial?
- How do you handle flaky tests without masking defects?

4. How does asynchronous differ from synchronous? Draw the architecture.

Interview-ready explanation

In synchronous execution, the caller sends a request and waits for the response before proceeding. It is straightforward for dependent steps but ties up the caller while waiting. In asynchronous execution, the caller initiates work and continues; completion is communicated later through a callback, future/promise, event, queue or polling. Async improves throughput and decoupling for I/O-bound work, but introduces ordering, timeout, retry, duplicate-message and observability concerns.

Practical example



```
// JavaScript illustration
console.log("A");
setTimeout(() => console.log("C - callback"), 0);
console.log("B");

// Output: A, B, C - callback
```

Senior-level considerations

Clarify that asynchronous does not automatically mean parallel execution. It means the caller is not blocked waiting for completion; actual concurrency depends on runtime and architecture. In tests, prefer condition-based waits with bounded timeouts over fixed sleeps.

Likely follow-up questions

- Is JavaScript single-threaded?
- Does async always mean parallel?
- How do you test eventual consistency?
- How do you prevent duplicate event processing?

Part 2 — API & Mocking

5. How can you mock a payment API when endpoints are the same?

Interview-ready explanation

Match requests using more than the URL. If several payment scenarios use the same endpoint, distinguish them by HTTP method, request body fields, headers, query parameters, correlation ID, customer/account, payment instrument or an explicit scenario header. The mock should return scenario-specific responses: approved, declined, insufficient funds, timeout, malformed response or duplicate transaction. Keep fixtures deterministic and avoid relying on request order.

Practical example

```
// Playwright: same endpoint, different request payloads
await page.route("**/api/payments", async route => {
  const req = route.request();
  const body = req.postDataJSON();

  if (body.amount === 0) {
    return route.fulfill({ status: 400, json: { code: "INVALID_AMOUNT" } });
  }
  if (body.testCard === "DECLINED") {
    return route.fulfill({ status: 402, json: { code: "CARD_DECLINED" } });
  }
  return route.fulfill({ status: 200, json: { status: "APPROVED", id: "mock-123" } });
});
```

Senior-level considerations

Prefer contract-driven mocks and schema validation. Use a dedicated mock server or service virtualization for integration tests; browser-level interception is excellent for UI edge cases but does not prove the backend integration works.

Likely follow-up questions

- What if the body is encrypted or tokenised?
- How do you avoid mocks diverging from the real contract?
- Would you mock at browser, service or gateway level?

6. How do you introduce delay in mocks?

Interview-ready explanation

Introduce a controlled delay at the mock layer to simulate slow dependencies. The delay should be configurable and bounded, and the test should assert the product's expected loading, timeout, retry or cancellation behaviour. Avoid arbitrary sleeps in the test itself when the purpose is to simulate a slow server.

Practical example

```
// Playwright delayed mock
await page.route("**/api/orders", async route => {
  await new Promise(resolve => setTimeout(resolve, 2500));
  await route.fulfill({
    status: 200,
    contentType: "application/json",
    body: JSON.stringify({ orders: [] })
  });
});
```

Senior-level considerations

Use short deterministic delays for functional tests and dedicated fault-injection tests for long latency. Validate spinner visibility, disabled duplicate-submit controls, user messaging, cancellation and telemetry. Never make the whole suite slow just to test one timeout path.

Likely follow-up questions

- How do you test client timeout vs server timeout?
 - How do you simulate a response that never completes?
 - How do you make delay tests reliable in CI?
-

7. Can you mock events? If yes, how?

Interview-ready explanation

Yes. Events can be mocked at different boundaries. For UI tests, emit a browser CustomEvent or intercept the network transport. For backend tests, publish a synthetic message to a test topic/queue or use a fake event producer. For consumer tests, feed a known event payload and verify state changes, idempotency and failure handling. The event contract should include event ID, type, timestamp, schema version and correlation/trace ID.

Practical example

```
// Browser event example
await page.evaluate(() => {
  window.dispatchEvent(new CustomEvent("payment-status", {
    detail: { eventId: "evt-101", status: "APPROVED" }
  }));
});

// Example event payload:
{
  "eventId": "evt-101",
  "type": "PaymentApproved",
  "schemaVersion": 1,
  "correlationId": "order-789",
  "data": { "orderId": "order-789", "amount": 125.50 }
}
```

Senior-level considerations

Test at-least-once delivery assumptions, idempotent consumers, schema compatibility, poison messages, retries and dead-letter handling. A mocked event proves consumer behaviour under controlled conditions; an end-to-end test should still validate the real broker integration.

Likely follow-up questions

- How do you test out-of-order events?
 - How do you test duplicate delivery?
 - How do you validate dead-letter queue handling?
-

8. How do you validate data retrieved by an API is correct?

Interview-ready explanation

Validate correctness at several layers: HTTP status and headers; response schema and data types; business rules; field-level values against an independent source of truth; relationships and referential integrity; persistence; and negative/edge cases. Do not merely assert that a field exists. For a read API, create or identify known data, call the API, and compare returned values with the expected record or database/service result using authorised test access.

Practical example

```
// REST Assured / Java illustration
given()
  .header("Authorization", "Bearer " + token)
  .when()
    .get("/api/members/" + memberId)
  .then()
    .statusCode(200)
    .body("memberId", equalTo(memberId))
    .body("status", equalTo("ACTIVE"))
    .body("age", greaterThanOrEqualTo(0))
    .body("effectiveDate", notNullValue());
```

Senior-level considerations

Use independent expected data; avoid deriving expected values from the same API response. Mask sensitive data in logs. For healthcare or financial domains, validate authorisation boundaries, tenant/state scoping, effective dates, audit trails and data minimisation.

Likely follow-up questions

- How do you validate a nested JSON response?
- How do you compare API data with DB data?
- How do you handle eventually consistent reads?

9. What tests do you conduct for an API?

Interview-ready explanation

Cover functional and non-functional dimensions: positive, negative, boundary and equivalence cases; schema/contract; authentication and authorisation; validation and error handling; HTTP method semantics; headers and content negotiation; pagination, sorting and filtering; idempotency; concurrency; rate limiting; backward compatibility; data integrity; performance; resilience; and security. Prioritise by risk and consumer impact.

Practical example

Category	Example
Positive	Valid request returns expected status and body
Negative	Missing field, invalid enum, malformed JSON
Boundary	Minimum/maximum values, date edges
Security	Unauthenticated, wrong role, IDOR
Contract	Schema, required fields, enum compatibility
Reliability	Retry, timeout, duplicate request, 5xx
Performance	Latency percentiles, throughput, error rate

Senior-level considerations

Use a test pyramid: many fast contract/unit/service tests, focused integration tests, and fewer end-to-end journeys. Keep test data isolated and ensure negative tests do not mutate production-like shared data.

Likely follow-up questions

- Which tests belong in CI on every PR?
 - How do you test backward compatibility?
 - How do you prioritise API coverage?
-

10. What is your metric for validating API performance?

Interview-ready explanation

Do not use average response time alone. Track latency percentiles (p50, p90, p95, p99), throughput (requests/second), error rate, timeout rate, saturation (CPU, memory, connection pools, thread pools), and resource utilisation. Include payload size and dependency timings where possible. Define pass/fail thresholds from business SLOs and workload expectations, not arbitrary universal numbers.

Practical example

Metric	What it tells you
p95 / p99 latency	Experience of slower requests / tail latency
Throughput	Requests handled per second at target load
Error rate	Failed requests / total requests
Apdex (if used)	User satisfaction approximation
CPU / memory	Resource pressure and saturation
DB pool / queue depth	Backend bottlenecks and backlog

Senior-level considerations

Report latency by endpoint and transaction, not only as one blended number. Include warm-up, ramp-up, steady state, spike and soak phases. Correlate client metrics with server traces and infrastructure metrics.

Likely follow-up questions

- Why can average latency hide a problem?
 - How do you distinguish load-generator saturation from app saturation?
 - How do you define a realistic workload?
-

Part 3 — Performance & Security

11. Why is JMeter used, and how is k6 different or better?

Interview-ready explanation

Apache JMeter is a mature, GUI-friendly, Java-based load-testing tool with a broad protocol ecosystem, extensive plugins and a large user base. Grafana k6 is code-first and developer-oriented: tests are authored in JavaScript, can be version-controlled, are easy to run in CI, and provide built-in metrics and threshold-based pass/fail criteria. It is not universally “better”. Choose based on protocol needs, team skills, ecosystem, scale model, reporting and operational constraints. Both can produce useful load tests when designed correctly.

Practical example

Dimension	JMeter	k6
Authoring	GUI test plans plus code options	Code-first JavaScript
CI/CD	CLI/non-GUI execution supported	Designed for automation and CI workflows
Ecosystem	Mature plugins and protocol support	Extensions; strong HTTP/API workflow
Team fit	Teams preferring GUI and broad plugin ecosystem	Teams preferring code review and Git
Decision	Choose for required protocol/plugin coverage	Choose for scriptable, threshold-driven workflows

```
// k6 example: simple API load test
import http from 'k6/http';
import { check } from 'k6';

export const options = {
  vus: 20,
  duration: '1m',
  thresholds: {
    http_req_failed: ['rate<0.01'],
    http_req_duration: ['p(95)<500'],
  },
};

export default function () {
  const res = http.get('https://test.example.com/api/products');
  check(res, { 'status is 200': r => r.status === 200 });
}
```

Senior-level considerations

Do not claim k6 is faster or better in every situation. Compare the actual protocols and extensions required. Validate the load generator itself, use production-like data and protect shared environments. k6 thresholds can make SLO expectations executable.

Likely follow-up questions

- What is the difference between load, stress, spike and soak testing?
- How do you calculate virtual users and arrival rate?
- Why can a load test pass while real users see slowness?

12. What can you do for security scanning of an application?

Interview-ready explanation

Use a layered security-testing strategy integrated into the SDLC. SAST analyses source code; SCA checks third-party dependencies and licences; secrets scanning detects exposed credentials; IaC/container scanning checks infrastructure definitions and images; DAST probes a running application; API security testing checks authentication, authorisation, input validation and abuse controls; and penetration testing explores attack paths. Map coverage to OWASP guidance and the organisation's threat model.

Practical example

Layer	Examples / purpose
SAST	Static code analysis: SonarQube, Semgrep, CodeQL
SCA	Dependency vulnerabilities: OWASP Dependency-Check, Snyk
DAST	Running app scan: OWASP ZAP, Burp Suite
Secrets	Detect keys/tokens committed to Git
Container / IaC	Scan images and Terraform/Kubernetes manifests
Manual security	Threat modelling, access-control tests, penetration testing

Senior-level considerations

Security testing requires permission and safe scope. Do not run destructive scans against production without explicit approval. Prioritise exploitability, exposure and business impact; assign owners and SLAs. A scanner is not a substitute for threat modelling or manual verification.

Likely follow-up questions

- How do you test IDOR/BOLA?
 - How do you manage false positives?
 - How do you scan authenticated workflows?
-


Part 4 — Observability & Leadership

13. How do you maintain test execution dashboards?

Interview-ready explanation

Build dashboards around decisions, not decoration. Ingest structured test results from CI (JUnit XML, Allure, Playwright JSON or equivalent) into a central store. Add metadata such as build, branch, commit, environment, suite, browser, state/tenant, owner and run ID. Display trend views for pass/fail/skip, first-run pass rate, flaky rate, duration, failure categories, coverage and release status. Link each failed test to logs, screenshots, video, trace and pipeline run.

Practical example



Useful dashboard panels: pass rate trend; failures by component; flaky tests; duration p50/p95; failures by environment/browser; top recurring defects; release gate status; last successful run; and links to evidence.

Senior-level considerations

Use a unique execution ID and idempotent ingestion. Separate test failures from infrastructure failures and product defects. Keep historical data long enough to identify trends. Never report retry-pass as an uncomplicated pass; show first attempt and final outcome distinctly.

Likely follow-up questions

- How do you prevent duplicate result ingestion?
- How do you define flaky rate?
- How do you handle reruns in the dashboard?

14. How do you process data for test execution?

Interview-ready explanation

Treat test data as a managed lifecycle. Identify required entities and dependencies, create data through supported APIs or fixtures, isolate each test/run with unique identifiers, seed reference data, protect sensitive information through masking or synthetic data, and clean up reliably. For parallel tests, avoid shared mutable records. Keep datasets versioned and environment-aware; record provenance so failures can be reproduced.

Practical example

```
# Example test-data naming convention
run_id = "build-4821-worker-03"
member_id = "qa-" + run_id
# Create member through a supported test API.
# Store returned IDs in the test context.
# Delete or expire the created records during teardown.
```

Senior-level considerations

Use synthetic data for sensitive domains. For healthcare insurance, ensure test records represent controlled states, coverage periods, household relationships and eligibility scenarios without exposing real member data. Make test data repeatable and auditable.

Likely follow-up questions

- When would you use database seeding vs API setup?
- How do you handle data that must be unique?
- How do you test date-sensitive eligibility rules?

15. What KPIs do you consider while testing?

Interview-ready explanation

Use a balanced set of quality, delivery and automation-health indicators. Examples: escaped defects, defect severity, requirement/risk coverage, test execution progress, first-run pass rate, flaky-test rate, automation coverage of critical journeys, mean time to diagnose, regression duration, API p95 latency, defect reopen rate and release-blocking defects. Avoid vanity metrics such as raw number of automated tests without risk context.

Practical example

KPI	How to interpret it
Escaped defects	Issues found after release; review severity and detection gap
Risk coverage	Critical business risks with effective tests
First-run pass rate	Reliability before retries
Flaky rate	Tests with inconsistent outcomes
Regression duration	Time to obtain actionable release confidence
MTTR / diagnosis time	Time to identify and resolve failures
Performance SLOs	Latency, throughput and error budgets

Senior-level considerations

Define each KPI precisely, its denominator, time window and owner. Segment by product area and risk. Pair quantitative metrics with qualitative release-risk analysis. Metrics should trigger investigation, not become individual performance targets in isolation.

Likely follow-up questions


- Can 100% automation coverage mean high quality?
- How do you avoid teams gaming KPIs?
- Which KPIs would you show executives vs engineers?

16. Can you explain the services involved in the product you test?

Interview-ready explanation

Explain the product as a business journey and then map the services. For a healthcare insurance exchange, a typical conceptual flow may include identity/authentication, consumer profile, household/application, eligibility rules, plan catalogue, pricing/premium calculation, enrolment, payment, document generation, notifications and downstream carrier/partner integrations. The exact service boundaries depend on the product architecture; never invent service names. Explain the data ownership, APIs/events, dependencies, failure modes and observability for the actual system.

Practical example



How to explain one service: purpose → input/output contract → data store/owner → upstream/downstream dependencies → security → failure handling → tests and monitoring.

Senior-level considerations

Prepare a truthful architecture story from your own product: draw the actual components, identify which are synchronous APIs and which are events, describe a real defect or failure scenario, and explain how your tests detect it. Clearly label any simplified or hypothetical architecture.

Likely follow-up questions

- What happens if eligibility succeeds but plan catalogue fails?
- How do you test state-specific rules?
- How do you trace one transaction across services?

17. An automation tester refuses to do manual testing. How do you handle it?

Interview-ready explanation

First understand the reason: unclear role expectations, lack of confidence, workload, accessibility or a belief that manual testing is beneath the role. Reconfirm the team's quality ownership and explain that automation is one technique, not a substitute for exploratory testing, risk analysis or acceptance validation. Agree on a balanced plan: the engineer can automate repeatable regression while also performing exploratory testing, reviewing requirements, validating new features and investigating failures. Coach, pair and set measurable expectations. Escalate only if the behaviour continues after expectations and support are clear.

Senior-level considerations

Lead with empathy and clarity. Do not equate manual testing with low skill. A strong QA engineer chooses the right technique for the risk. Make responsibilities explicit across the team and protect time for exploratory testing.

Likely follow-up questions

- How do you handle a senior engineer who disagrees?
- What if sprint commitments leave no time for exploratory testing?
- How do you avoid using manual testing as a punishment?

Part 5 — Testing Strategy

18. How can you test Add to Cart using both UI and API?

Interview-ready explanation

Split the journey by test objective. API tests validate cart service rules quickly: add item, quantity, price, inventory, duplicate item behaviour, promotions, tax/shipping calculations, persistence and error cases. UI tests validate that a user can search/select a product, add it, see the cart count and line item, update quantity, remove it and receive correct messages. Use API setup to create preconditions, but keep a small number of end-to-end UI tests that exercise the real user flow. Cross-check UI state against the API response where useful.

Practical example

```
// Playwright example: UI journey + API assertion (illustrative)
await page.goto('/products/sku-123');
await page.getByRole('button', { name: 'Add to cart' }).click();
await expect(page.getByTestId('cart-count')).toHaveText('1');

const response = await request.get('/api/cart');
expect(response.ok()).toBeTruthy();

const cart = await response.json();
expect(cart.items).toHaveLength(1);
expect(cart.items[0].sku).toBe('sku-123');
expect(cart.items[0].quantity).toBe(1);
```

Senior-level considerations

Test guest and authenticated users, session persistence, multiple tabs/devices, stock changes, discounts, currency, rounding, stale prices, idempotency and checkout hand-off. Do not assert only the cart badge; validate the underlying business state.

Likely follow-up questions

- How do you test price changes between product page and checkout?
- How do you test inventory race conditions?
- How do you avoid test pollution in a shared cart?

19. How can you create a Page Object for the Amazon homepage?

Interview-ready explanation

Model the page around user-visible responsibilities, not every DOM element. Keep locators and page actions in a page object; keep business assertions in tests or domain-level helpers. Use accessible roles/labels where possible and stable test IDs when available. Amazon's real site is dynamic and may vary by region, consent, experiments and authentication, so for an interview use a simplified illustrative page object rather than assuming exact live selectors.

Practical example

```
import { Page, Locator } from '@playwright/test';

export class HomePage {
  readonly searchBox: Locator;
  readonly searchButton: Locator;
```

```
readonly cartLink: Locator;

constructor(private page: Page) {
  this.searchBox = page.getByRole('searchbox');
  this.searchButton = page.getByRole('button', { name: /search/i });
  this.cartLink = page.getByRole('link', { name: /cart/i });
}

async open() {
  await this.page.goto('/');
}

async searchProduct(term: string) {
  await this.searchBox.fill(term);
  await this.searchButton.click();
}

async openCart() {
  await this.cartLink.click();
}
}
```

Senior-level considerations

Explain the boundary: page object encapsulates UI mechanics and reusable actions; test describes scenario and verifies business outcome. Avoid giant page objects, generic click wrappers and hidden waits. For a search flow, consider a separate SearchResultsPage object.

Likely follow-up questions

- Would you put expect() assertions inside the page object?
 - How would you handle a location/region pop-up?
 - How do you avoid brittle selectors?
-

Part 6 — Coding / DSA

20. Find even and odd digits in 87530 (classify each digit).

Interview-ready explanation

Convert the number to text or repeatedly use modulo and division. For a simple interview answer, a string is readable and preserves each digit. For numeric processing, use modulo 10 to extract the last digit and divide by 10. Note that the number 87530 contains even digits 8, 0 and odd digits 7, 5, 3.

Practical example

```
public class EvenOddDigits {
    public static void main(String[] args) {
        String number = "87530";

        for (int i = 0; i < number.length(); i++) {
            int digit = number.charAt(i) - '0';

            if (digit % 2 == 0) {
                System.out.println(digit + " is even");
            } else {
                System.out.println(digit + " is odd");
            }
        }
    }
}

// Output:
// 8 is even
// 7 is odd
// 5 is odd
// 3 is odd
// 0 is even
```

Senior-level considerations

Complexity: $O(d)$ time for d digits and $O(1)$ extra space, excluding the input string representation. A numeric modulo solution also takes $O(d)$ time and $O(1)$ extra space. Clarify whether leading zeroes matter; numeric input cannot preserve them.

Likely follow-up questions

- How would you count even and odd digits?
- How would you handle negative numbers?
- Can you solve it without converting to String?

21. Find the maximum binary gap for number 529.

Interview-ready explanation

A binary gap is a maximal sequence of consecutive zeroes that is surrounded by ones at both ends in the binary representation. Convert 529 to binary: $529 = 512 + 16 + 1$, so binary is 1000010001. The zero runs between ones have lengths 4 and 3. The maximum binary gap is therefore 4. Trailing zeroes do not count because they are not followed by a one.

Practical example

```
public class BinaryGap {
    public static int maxBinaryGap(int n) {
```

```

String binary = Integer.toBinaryString(n);
int maxGap = 0;
int currentGap = 0;
boolean seenOne = false;

for (char bit : binary.toCharArray()) {
    if (bit == '1') {
        if (seenOne) {
            maxGap = Math.max(maxGap, currentGap);
        }
        seenOne = true;
        currentGap = 0;
    } else if (seenOne) {
        currentGap++;
    }
}
return maxGap;
}

public static void main(String[] args) {
    System.out.println(maxBinaryGap(529)); // 4
}

```

Senior-level considerations

Complexity: $O(\log n)$ time and $O(\log n)$ space for the string approach. Test edge cases: 1 (gap 0), 9 (1001 → gap 2), 32 (100000 → gap 0), and 529 (gap 4).

Likely follow-up questions

- What is the binary representation of 20?
- Why don't trailing zeroes count?
- Can you solve it using bit operations instead of a String?

Part 7 — How to handle cross-examination

Interviewers often use follow-ups to check whether the candidate understands the limits of a solution. Prepare to answer these consistently:

- Why did you choose this design? State the requirement and constraint that drove the choice.
- What would fail first at 5× the current scale? Identify the bottleneck and how you would measure it.
- How do you know the test itself is reliable? Discuss isolation, determinism, evidence and first-run results.
- What happens when a dependency is slow or unavailable? Explain timeout, retry, fallback, circuit breaker and user-visible behaviour where relevant.
- How would you roll this out safely? Start small, measure, compare, monitor and have a rollback path.
- What did you personally own? Separate your direct contribution from the team's work.
- What would you improve now? Show judgement without dismissing the existing design.

A 60-second structure for answering any senior question

Time	What to say
0–10 sec	Restate the problem and assumptions.
10–25 sec	Describe the architecture or approach.
25–40 sec	Give a concrete example and test coverage.
40–50 sec	Explain trade-offs and failure handling.

Time	What to say
50–60 sec	Name metrics and how you know it works.

Claims to avoid unless you can defend them

- “Our framework is 100% automation.” Explain what is covered, what is not, and why.
- “We eliminated flakiness.” Prefer measurable first-run reliability trends and remaining known causes.
- “k6 is always better than JMeter.” Explain the workload, protocol and team context.
- “Grid is fully asynchronous.” Explain the synchronous HTTP and asynchronous Event Bus paths.
- “We test security with a scanner.” Explain threat model, access-control tests, manual verification and remediation.

Quick revision sheet

Topic	Must-remember points
Selenium Grid 4	Router, New Session Queue, Distributor, Node, Session Map, Event Bus; HTTP + async events.
Framework design	Layering, isolation, fixtures, data, config, reporting, observability, maintainability.
Parallelism	Independent sessions/contexts, unique data, bounded workers, reliable teardown.
Async vs sync	Wait-for-response vs initiate-and-continue; handle ordering, retries and timeouts.
Mocking	Match method + path + payload/headers; contract-based fixtures; simulate latency/events.
API validation	Status, schema, business rules, source-of-truth, security, negative cases.
API performance	p95/p99, throughput, error rate, saturation, realistic workload and SLOs.
Performance tools	JMeter mature GUI/plugin ecosystem; k6 code-first scripts and thresholds.
Security	SAST, SCA, DAST, secrets, IaC/container, manual threat-based tests.
Dashboards	Unique run ID, trends, flaky/first-run rate, duration, failure evidence.
Leadership	Empathy, clear expectations, coaching, exploratory testing and shared quality.
UI + API	API for breadth and setup; UI for user journey; verify business state.
DSA	87530 → even 8,0; odd 7,5,3. 529 → 1000010001 → max gap 4.

Official references for version-sensitive details

- Selenium Grid architecture: <https://www.selenium.dev/documentation/grid/architecture/>
- Selenium Grid components: <https://www.selenium.dev/documentation/grid/components/>
- Selenium Grid getting started: https://www.selenium.dev/documentation/grid/getting_started/
- Playwright network mocking: <https://playwright.dev/docs/network>
- Playwright Route API: <https://playwright.dev/docs/api/class-route>
- Grafana k6 metrics: <https://grafana.com/docs/k6/latest/using-k6/metrics/>
- Grafana k6 thresholds: <https://grafana.com/docs/k6/latest/using-k6/thresholds/>
- OWASP Web Security Testing Guide: <https://owasp.org/www-project-web-security-testing-guide/>
- OWASP ZAP documentation: <https://www.zaproxy.org/docs/>
- Apache JMeter user manual: <https://jmeter.apache.org/usermanual/>

Final reminder: explain the problem, show the design, justify the trade-offs, prove the result with evidence, and be honest about what you personally implemented.